

Job Description Data Engineering

Campus Hiring-2027



About the Company:

Celebal Technologies is a software services company working in Data Science, Big Data, and Enterprise Cloud. We help large organisations make better use of their data through well-built pipelines, clean architectures, and reliable cloud infrastructure. Our engineering teams work with tools like Databricks, Apache Spark, Azure Data Factory, and Python on real client projects.

We are a Microsoft Gold Partner and work mostly within the Azure ecosystem. Our clients include Fortune 500 companies across banking, healthcare, manufacturing, retail, and energy.

About the Role:

We are hiring fresh graduates as Data Engineers. You will learn to build and maintain data pipelines, work with cloud platforms, and write code that moves and transforms data for real business use cases. You will be trained on the job and mentored by senior engineers, but we do expect you to come in with a solid foundation in programming, SQL, and basic understanding of how databases and distributed systems work.

Over the course of your first year, you will pick up PySpark, get comfortable with Databricks, learn how cloud services like Azure Data Factory and Azure Storage work, and start contributing to client projects. As you grow, you will also get exposure to building data infrastructure for AI and ML use cases, including connecting data systems to language models.

This is a hands-on coding role. We are looking for people who enjoy writing code, solving data problems, and building things that work.

Eligibility Criteria

- B.Tech or M.Tech in Computer Science, AI/ML, IT, Data Science, or similar
- No active backlogs

- Passing out year: 2027
- Minimum 7.0 CGPA (or 65%) throughout
- Should have done coursework in data structures, algorithms, databases, or operating systems
- At least one project involving data handling — could be a data pipeline, a database project, a web scraping tool, a cloud deployment, or anything where you moved and processed data programmatically



Designation

Data Engineer

Key Responsibility Areas (KRA's)

- Learn and work with PySpark and Databricks to build data pipelines that run on Azure or AWS.
- Write Python and SQL code for data extraction, transformation, and loading (ETL) tasks.
- Work with Databricks features like Delta Lake tables, Unity Catalog, job scheduling, and notebook-based development.
- Help set up and maintain ETL workflows using Azure Data Factory, Databricks Workflows, or Airflow.
- Connect to and pull data from different sources: databases (SQL Server, PostgreSQL, MongoDB), APIs, file systems, and streaming platforms like Kafka or Event Hubs.
- Learn data modelling basics: how to design tables using star schemas, handle slowly changing data, and organise data for downstream use.
- Support the team in deploying pipelines using Docker, Git, and CI/CD tools like Azure DevOps or GitHub Actions.
- Get exposure to building data infrastructure for AI/ML use cases — feeding data into models, setting up training jobs, or connecting data sources to language models through MCP servers.
- Follow data governance practices: access controls, data quality checks, and proper documentation.
- Collaborate with senior engineers, data scientists, and analysts to understand data requirements and deliver clean, reliable datasets.
- Write clear documentation and maintain your code in version control.

Desired Skills

- Python — comfortable writing scripts, not just running notebooks
- SQL — can write joins, subqueries, window functions, and understands indexing basics
- Understanding of data structures and algorithms
- Basic knowledge of how databases work (relational and non-relational)
- Exposure to any cloud platform: Azure, AWS, or GCP — even if through a course or personal project
- Familiarity with PySpark or willingness to learn it quickly
- Awareness of what Databricks is and how lakehouse architecture works (even from reading or a tutorial)
- Git basics: clone, commit, push, pull, branching
- Comfortable working in Linux/terminal
- Basic understanding of how ETL pipelines work — what happens between raw data and a clean table
- Good problem-solving habits — can break a problem into steps and code a solution
- Willingness to learn on the job — you will not know everything on day one, and that is fine

Good to Have

- Hands-on experience with PySpark or Spark (even from a course project or Databricks community edition)
- Any project on Azure or AWS — even a small deployment counts
- Experience with Docker or any containerisation tool
- Built a data pipeline end to end: ingestion, transformation, storage, and some form of output
- Exposure to streaming: Kafka, Spark Streaming, or Event Hubs
- Used Airflow, dbt, or any workflow orchestration tool
- Built or worked on an AI-related project — could be an MCP server, a model pipeline, an LLM integration, or a simple ML deployment
- Basic understanding of how ML models consume data
- Terraform or any Infrastructure as Code experience
- Scala or Java (even basic level)
- Open-source contributions or personal projects on GitHub